

```
In [1]: import pandas as pd
import numpy as np

In [2]: from tensorflow.keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 6S 8us/step
```

Visualizing the Image Data

[illegible]

```
Out[9]: array[[5, 0, 4, ..., 5, 6, 8], dtype=uint8]

In [10]: y_test
Out[10]: array[[7, 2, 1, ..., 4, 5, 6], dtype=uint8]

In [11]: from tensorflow.keras.utils import to_categorical

In [12]: y_train.shape
Out[12]: (60000,)
```

```
In [13]: y_example = to_categorical(y_train)

In [14]: y_example
Out[14]: array([[0., 0., 0., ..., 0., 0., 0.],
 [1., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [9., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)

In [15]: y_example.shape
Out[15]: (60000, 10)

In [16]: y_example[0]
Out[16]: array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.], dtype=float32)

In [17]: y_cat_test = to_categorical(y_test, 10)

In [18]: y_cat_train = to_categorical(y_train, 10)
```

Processing X Data

```
In [19]: single_image.max()
```

```
Out[20]: 0
```

```
In [21]: x_train = x_train / 255
         x_test = x_test / 255
```

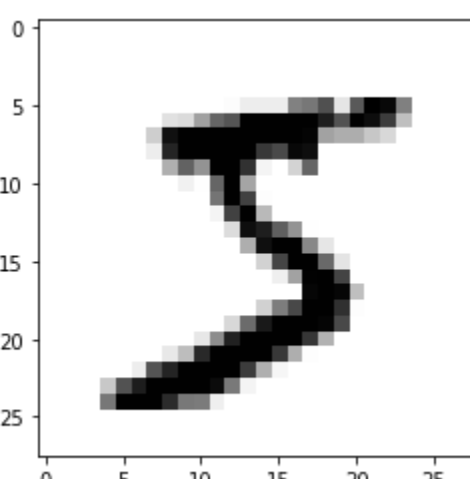
```
In [22]: scaled_single = x_train[0]
```

```
In [23]: scaled_single.max()
```

```
Out[23]: 1.0
```

```
In [24]: plt.imshow(scaled_single, cmap='gray_r')
```

```
Out[24]: <matplotlib.image.AxesImage at 0x7f8c7cbf62d0>
```



A grayscale plot of a handwritten digit '5' on a 28x28 pixel grid. The digit is rendered in black on a white background. The plot includes x and y axes ranging from 0 to 25.

Reshaping the Data

```
In [25]: x_train.shape
```

```
Out[25]: (60000, 28, 28)
```

```
In [26]: x_test.shape
```

```
Out[26]: (10000, 28, 28)
```

```
x_train = x_train.reshape
```

```
In [20]: x_train.shape
Out[20]: (60000, 28, 28, 1)

In [29]: x_test = x_test.reshape(10000, 28, 28, 1)

In [30]: x_test.shape
Out[30]: (10000, 28, 28, 1)
```

Training the Model

```
In [31]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D, Flatten

In [32]: model = Sequential()

# CONVOLUTIONAL LAYER
model.add(Conv2D(filters=32, kernel_size=(4, 4), input_shape=(28, 28, 1), activation='relu'))

# POOLING LAYER
model.add(MaxPool2D(pool_size=(2, 2)))

# FLATTEN IMAGES FROM 28 by 28 to 764 BEFORE FINAL LAYER
```

```
model.add(Flatten())
```

```

model.add(Dense(128, activation='relu'))

# LAST LAYER IS THE CLASSIFIER, THUS 10 POSSIBLE CLASSES
model.add(Dense(10, activation='softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

In [33]: model.summary()

Model: "sequential"

Layer (type)                Output Shape              Param #
=====
conv2d (Conv2D)              (None, 25, 25, 32)        544
max_pooling2d (MaxPooling2D) (None, 12, 12, 32)         0
flatten (Flatten)            (None, 4608)               0
dense (Dense)                (None, 128)                589952
dense_1 (Dense)              (None, 10)                 1290
=====
Total params: 591,786
Trainable params: 591,786
Non-trainable params: 0

In [34]: from tensorflow.keras.callbacks import EarlyStopping

In [35]: early_stop = EarlyStopping(monitor='val_loss', patience=2)

```

```

y_cat_train, epochs=10,
validation_data=(x_test, y_cat_test),
callbacks=[early_stop])

Epoch 1/10
1875/1875 [=====] - 19s 5ms/step - loss: 0.1384 - accuracy: 0.9582 - val_loss: 0.0571 - val_
accuracy: 0.9620
Epoch 2/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.0468 - accuracy: 0.9860 - val_loss: 0.0426 - val_a
ccuracy: 0.9863
Epoch 3/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.0399 - accuracy: 0.9904 - val_loss: 0.0467 - val_a
ccuracy: 0.9854
Epoch 4/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.0209 - accuracy: 0.9931 - val_loss: 0.0432 - val_a
ccuracy: 0.9880

Out[36]: <keras.callbacks.History at 0x7f8c165cd1d0>

```

```
model.metrics_names
```

```
In [38]: losses = pd.DataFrame(model.history.history)
```

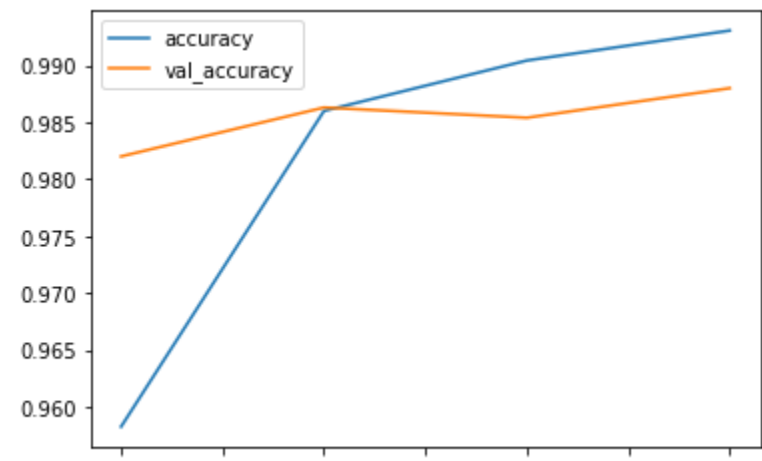
```
In [39]: losses.head()
```

Out[39]:

	loss	accuracy	val_loss	val_accuracy
0	0.138413	0.958233	0.057086	0.9820
1	0.046792	0.986000	0.042635	0.9863
2	0.030693	0.990433	0.046735	0.9854
3	0.020906	0.993067	0.043235	0.9880

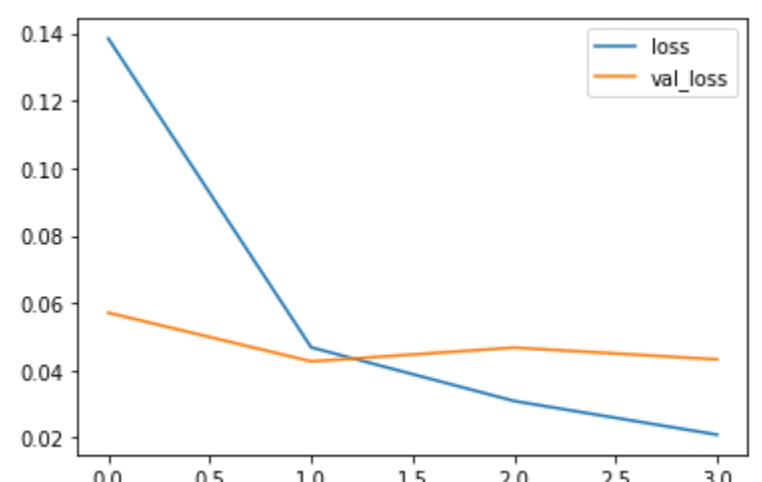
```
In [40]: losses[['accuracy', 'val_accuracy']].plot()
```

Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8c162a9890>



```
In [41]: losses[['loss', 'val_loss']].plot()
```

Out[41]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8c161e3310>



```
In [42]: print(model.metrics_names)
print(model.evaluate(x_test, y_cat_test, verbose=0))
```

['loss', 'accuracy']
0.02090607568132401 0.98799997568132401

```
In [43]: from sklearn.metrics import classification_report, confusion_matrix

In [44]: # predictions = model.predict_classes(x_test)
predictions = (model.predict(x_test) > 0.5).argmax(axis=1).astype('uint8')
313/313 [=====] - 1s 2ms/step

In [45]: y_cat_test.shape
Out[45]: (10000, 10)

In [46]: y_cat_test[0]
Out[46]: array([0., 0., 0., 0., 0., 0., 0., 1., 0., 0.], dtype=float32)

In [47]: predictions[0]
Out[47]: 7
```

```
In [48]: y_test
Out[48]: array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)

In [49]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	980
1	1.00	0.99	1.00	1135
2	0.97	1.00	0.98	1032
3	0.99	0.99	0.99	1010
4	0.99	1.00	0.99	902
5	0.99	0.99	0.99	892

```
6      1.00      0.98      0.99      958
7      0.99      0.98      0.98      1028
8      0.98      0.98      0.98      974
9      0.99      0.97      0.98      1009

accuracy      0.99      0.99      0.99      10000
macro avg     0.99      0.99      0.99      10000
weighted avg   0.99      0.99      0.99      10000
```

```
In [50]: confusion_matrix(y_test, predictions)
```

```
Out[50]: array([[ 976,    0,    1,    0,    0,    0,    1,    0,    1,    1],
                 [    0,  1129,    4,    0,    0,    1,    0,    0,    1,    0],
                 [    0,    0,  1030,    0,    0,    0,    0,    2,    0,    0],
                 [    0,    0,    0,   998,    0,    0,    0,    0,    0,    0],
                 [    0,    0,    0,    0,   998,    0,    0,    0,    0,    0],
                 [    0,    0,    0,    0,    0,   998,    0,    0,    0,    0],
                 [    0,    0,    0,    0,    0,    0,   998,    0,    0,    0],
                 [    0,    0,    0,    0,    0,    0,    0,   998,    0,    0],
                 [    0,    0,    0,    0,    0,    0,    0,    0,   998,    0],
                 [    0,    0,    0,    0,    0,    0,    0,    0,    0,   998]])
```

```
[ [ 1, 0, 0, 0, 0, 978, 0, 0, 0, 0, 3],
  [ 2, 0, 0, 0, 6, 0, 881, 0, 0, 2, 0],
  [ 4, 3, 0, 0, 0, 1, 0, 942, 0, 4, 0],
  [ 3, 0, 15, 0, 0, 0, 0, 1086, 2, 2],
  [ 0, 0, 5, 1, 0, 0, 0, 2, 958, 2],
  [ 6, 2, 3, 2, 0, 3, 0, 3, 2, 970]]])
```

```
In [51]: import seaborn as sns

In [52]: plt.figure(figsize=(10, 6))
sns.heatmap(confusion_matrix(y_test, predictions), annot=True)

Out[52]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8be06de9e0>
```

[illegible]

```

In [53]: my_number = x_test[0]

In [54]: plt.imshow(my_number.reshape(28, 28), cmap='gray_r')

Out[54]: <matplotlib.image.AxesImage at 0x7f8bfc15b960>

```